


An aerial photograph of a city, likely London, showing a dense grid of streets and buildings. A river, possibly the River Tyne, runs vertically through the center of the image. The buildings are mostly rectangular and arranged in a regular pattern. The river is a dark, straight line. The overall color palette is muted, with greys, browns, and greens.

Evaluation & Audits

Ethics and Bias in AI – lecture 6

Dr Shauna Concannon

Shauna.j.Concannon@durham.ac.uk

Recap



So far we have explored:

- Bias, fairness & discrimination –definitions
 - Impact of bias on individuals & groups
 - Introduction to fairness in ML (e.g. fairness through unawareness, FPP, FNP, etc.)
 - Design frameworks for anticipating and minimising harms (e.g. VSD)
-
- Next: look at some applied examples of bias identification and mitigation efforts
 - Consider contexts where categorical demographic information isn't readily available...

This week: Evaluation and Audits


- Bias in machine vision & facial processing technologies
 - (Representational bias & harm)
- Audits, Systematic reviews of outputs
- Benchmarks, domain specific bias metrics, evaluation measures
- Critiques & considerations

To consider:

- How do we practically identify bias?
- What if we only have access to outputs?
- How do we create meaningful metrics?
- What if demographic data isn't available?


Bias in facial processing technologies

- Facial recognition systems
 - Deployed in contexts such as policing despite known issues (e.g. Gender Shades project)
- Image annotation (ImageNet)
 - over 14 million labeled images organized into more than 20 thousand categories.
- Image generation (Stable Diffusion, DALL-E)



Bias in facial processing technologies

- Facial recognition systems
 - Deployed in contexts such as policing despite known issues (e.g. Gender Shades project)
- Image annotation (ImageNet)
 - over 14 million labeled images organized into more than 20 thousand categories.
- Image generation (Stable Diffusion, DALL-E)




#ImageNetRoulette

- Art project by Kate Crawford and Trevor Paglen (2019)
- Read more:
<https://excavating.ai/>

No matter what kind of image I upload, ImageNet Roulette, which categorizes people based on an AI that knows 2500 tags, only sees me as Black, Black African, Negroid or Negro.

Some of the other possible tags, for example, are “Doctor,” “Parent” or “Handsome.”



ImageNet Roulette: An Experiment in Classification

The ImageNet dataset is typically used for object recognition. But as part of our archeological method, we were interested to see what would happen if we trained an AI model exclusively on its “person” categories. The result of that experiment is ImageNet Roulette.


ImageNet Roulette uses an open-source Caffe deep-learning framework (produced at UC Berkeley) trained on the images and labels in the “person” categories (which are currently “down for maintenance”). Proper nouns were removed.

When a user uploads a picture, the application first runs a face detector to find any faces. If it finds any, it sends them to the Caffe model for classification. The application then returns the original images with a bounding box showing the detected face and the label the classifier has assigned to the

An online image database will remove 600,000 pictures after an art project revealed the system's racist bias.

Christy Kuesel

Sep 24, 2019 5:20PM, via [The Art Newspaper](#)



Source: <https://www.artsy.net/article/artsy-editorial-online-image-database-will-remove-600-000-pictures-art-project-revealed-systems-racist-bias>

IMGENET

14,197,122 images, 21841 synsets index

[Home](#) [Download](#) [Challenges](#) [About](#)

Not logged in. [Log in](#)

Towards Fairer Datasets: Filtering and Balancing the Distribution of the People Subtree in the ImageNet Hierarchy

Kaiyu Yang (Princeton University), Klint Qinami (Princeton University), Li Fei-Fei (Stanford University), Jia Deng (Princeton University), Olga Russakovsky (Princeton University)

September 17, 2019

This research post describes the effort to improve ImageNet data, initiated, conceptualized, and executed solely by the above authors as part of a [NSF funded project](#) that started one year ago on September 1, 2018. The full technical report is published at [FAT* 2020](#).

[\[paper\]](#) [\[project\]](#)

As AI technology advances from research lab curiosities into people's daily lives, ensuring that AI systems produce appropriate and fair results has become an important scientific question. This requires multiple iterations of research on the data, algorithms, and design choices in all stages of building an AI system. To this end, ImageNet, as an influential research dataset, deserves to be critically examined. Below we summarize the most recent effort from our research team on this topic. In this work, we consider three key factors within ImageNet that may lead to problematic behavior in downstream computer vision technology: (1) the stagnant concept vocabulary, (2) the exhaustive illustration of all concepts with images, and (3) the inequality of representation in the images within concepts. We seek to illuminate the root causes of these concerns and take the first steps to mitigate them constructively. The full technical report was submitted for peer review in August 2019 and will be made available shortly.

Source: <https://image-net.org/update-sep-17-2019>

“Science progresses through trial and error, through understanding the limitations and flaws of past results ... We believe that ImageNet, as an influential research dataset, deserves to be critically examined, in order for the research community to design better collection methods and build better datasets.”

ImageNet statement

“There is no easy technical ‘fix’ by shifting demographics, deleting offensive terms, or seeking equal representation by skin tone,” they write. “The whole endeavor of collecting images, categorizing them, and labeling them is itself a form of politics, filled with questions about who gets to decide what images mean and what kinds of social and political work those representations perform.”

Paglen & Crawford

Source: <https://news.artnet.com/art-world/imagenet-roulette-trevor-paglen-kate-crawford-1658305>

Algorithmic auditing

‘When we specify that **public interest scrutiny of algorithms is important** we do not mean scrutiny that can be satisfied by legal tools alone, or scrutiny that can be achieved by a single instance of trial-and-error use of an algorithm. Although individual trial-and-error may detect some forms of harm (Edelman, 2010), **there are a variety of other problems that can only be uncovered via a systematic statistical investigation.**

Because algorithms often operate in a personalized way, **individual investigations are unlikely to produce a broad picture of the system’s operation across the diversity of its users’.**

Algorithmic audits

- An algorithmic audit is a **systematic** analysis of an algorithm or set of algorithms to identify any potential biases, errors, or other issues that may be present
- Typically carried out by a **team of experts** in the field who are well-versed in the technical aspects of algorithms, as well as the ethical and social implications of their use
- Can be conducted **by internal or by external** parties
- **Doesn't require access** to whole system necessarily, so useful for investigating commercial systems



Algorithmic audits

The goals of an algorithmic audit:

- Evaluate system performance wrt fairness, equality, ethics
- Identify system limitations and areas for improvement
- Ensure the algorithm / system is fair and equitable



Algorithmic audits

To design an algorithmic audit, you need to:

- Define the scope of the audit
 - including the specific algorithm or set of algorithms being evaluated, as well as the goals and objectives of the audit.
 - E.g. specifying the specific types of biases to be investigated
- Develop a methodology
 - Identify what criteria/metrics are used
 - Define a process for testing and evaluation
 - Establishing benchmarks and performance standards



Algorithmic audits

Outcomes & Actions:

- Demonstrate robust methodology & be explicit about the findings demonstrate
- Document findings and communicate to relevant stakeholders to highlight system limitations
- Encourage relevant parties to take action:
 - implement modifications
 - develop new policies and procedures to ensure ongoing accountability and transparency

Bias in facial processing technologies

- Gender Shades (Buolamwini & Gebru)
 - Domain: facial processing technologies
 - <https://youtu.be/TWWsW1w-BVo>

Gender Shades (Buolamwini & Gebru, 2018)


- Example of an external audit
 - Identify an appropriate dataset to test on
 - Generating necessary annotations in absence of demographic data
-
- First assessed the suitability of existing datasets for intersectional benchmarking
 - Provided skin type annotations for subjects in selected datasets using the dermatologist approved Fitzpatrick Skin Type classification system
 - Analysis of the IJB-A and Adience benchmarks revealed overrepresentation of lighter males, underrepresentation of darker females, and underrepresentation of darker individuals in general.



Gender Shades (Buolamwini & Gebru, 2018)


- Developed the Pilot Parliaments Benchmark (PPB) to achieve better intersectional representation on the basis of gender and skin type.
 - Images of parliamentarians, public figures with known identities and photos available under non-restrictive licenses posted on government websites.
- PPB consists of 1270 individuals from three African countries (Rwanda, Senegal, South Africa) and three European countries (Iceland, Finland, Sweden) selected for gender parity in the national parliaments.

Intersectionality



Pilot Parliaments Benchmark

Gender Shades (Buolamwini & Gebru, 2018)



Gender Shades (Buolamwini & Gebru, 2018)

- Evaluated performance of three commercial systems


Classifier	Metric	DF	DM	LF	LM
MSFT	PPV(%)	76.2	100	100	100
	Error Rate(%)	23.8	0.0	0.0	0.0
	TPR(%)	100	84.2	100	100
	FPR(%)	15.8	0.0	0.0	0.0
Face++	PPV(%)	64.0	99.5	100	100
	Error Rate(%)	36.0	0.5	0.0	0.0
	TPR(%)	99.0	77.8	100	96.9
	FPR(%)	22.2	1.03	3.08	0.0
IBM	PPV(%)	66.9	94.3	100	98.4
	Error Rate(%)	33.1	5.7	0.0	1.6
	TPR(%)	90.4	78.0	96.4	100
	FPR(%)	22.0	9.7	0.0	3.6



Internal vs external audits

- Buolamwini and Gebru also report on reactions from the companies whose services were audited to illustrate more and less effective responses
- External audits can put pressure on and create momentum inside of companies to respond to external critique. May then lead to concrete changes to their products.
- Internal auditing efforts have access to more data, allowing companies to conduct audits before public releases, establishing processes for companies, and allow them to take responsibility for their impact.

Types of AI audits



- Casper, Stephen, et al. "Black-box access is insufficient for rigorous ai audits." *The 2024 ACM Conference on Fairness, Accountability, and Transparency*. 2024.

Goal	Technique	Advantage
Identifying problems	Attacks	White-box methods are more reliable for detecting anomalous failures.
		White-box attacks are more efficient.
		White-box methods provide auditors with more attack options.
		Robustness to white-box attacks confers greater robustness assurances.
	Fine-tuning	Fine-tuning enables searching for harmful dormant capabilities.
	Interpretability	Interpretability tools enable stronger assurances of what knowledge and capabilities a system has.
Incentivizing responsible development	Outside-the-box assessment	Information helps auditors target complementary evaluations.
		Information about deployment assists with assessing societal impacts.
		Auditors can trace problems to datasets and methodological decisions made by developers.
Increasing transparency	Interpretability	Auditors can assess risk-mitigation strategies.
		White-box methods can address misconceptions about model outputs, characteristics, and risks.
		White-box methods assist with providing more accurate and meaningful explanations.
Enabling debugging	Attacks	Stronger attacks produce stronger instances of failures to incorporate into training data.
	Interpretability	Interpretability techniques enable more precise debugging.


Table 2: A summary of the advantages that various white-box and outside-the-box auditing techniques provide over black-box methods that are discussed in Section 4 and Section 5. In Appendix A, we expand on the variety of motivations for AI audits.

Casper, Stephen, et al. "Black-box access is insufficient for rigorous ai audits." *The 2024 ACM Conference on Fairness, Accountability, and Transparency*. 2024.



Key take aways

- The concept of the algorithmic audit aims to assess how algorithms contribute to social inequality,
- While they are often performed once to establish the presence of biases in various systems, should instead be approached as regular health checks (not a one-off process)
- Useful if you don't have access to underlying system, but even more so if you do
- But require careful scoping of metrics and test datasets – and how you communicate what an audit can / cannot demonstrate



Evaluation metrics & benchmarks

Evaluation metrics & benchmarks



Benchmarks are commonly used in ML

Datasets + standardised metrics



Aim to provide a consistent framework for evaluation



A number of specialist benchmark datasets and evaluation metrics have been introduced to address specific types of bias



Benchmarks
are becoming
increasingly
important in
AI Governance

EU AI Act Article 51

“A general-purpose AI model shall be classified as a general-purpose AI model with systemic risk if it meets any of the following conditions:

(a) it has high impact capabilities evaluated on the basis of appropriate technical tools and methodologies, including indicators and benchmarks [...]



Benchmarks are becoming increasingly important in AI Governance

UK AISI Inspect Framework

- Open-source framework for evaluating language models
- Facilitates “prompt engineering, tool usage, multi-turn dialog, and model graded evaluations”
- Features internally developed evaluations – and public benchmarks

Inspect is not yet being used as part of official regulatory efforts – but capacities are being ramped up across AISIs to fulfill the function of independent evaluators for planned and future regulation

Evaluation metrics & benchmarks

Specialist benchmark datasets and evaluation metrics attend to specific types of bias.

For example:

- Winogender
- WEAT (Word Embedding Association Test)
- Relative Negative Sentiment Bias (RNSB)

Example 1

A man and his son get into a terrible car crash. The father dies, and the boy is badly injured. In the hospital, the surgeon looks at the patient and exclaims, "I can't operate on this boy, he's my son!" How can this be?




Winogender

Gender Bias in Coreference Resolution (Rudinger et al., NAACL 2018)

Benchmark dataset that is used to evaluate the gender bias in NLP models.

- Co-reference resolution: the dataset consists of a set of sentences with a blank space where a pronoun would go, and the task is to predict the correct pronoun.
- The dataset includes examples where the correct pronoun could be either masculine or feminine, allowing for an evaluation of the model's ability to correctly predict gender-neutral pronouns.

Winogender: co-reference tasks



Winogender

Created an evaluation set hand-crafted sentences designed to test gender bias in co-reference resolution

1. **OCCUPATION** , a person referred to by their occupation and a definite article, e.g., “the paramedic.”
2. **PARTICIPANT** , a secondary (human) participant, e.g., “the passenger.”
3. **PRONOUN** , a pronoun that is coreferent with either OCCUPATION or PARTICIPANT.

Example:

The nurse notified the patient that...

1. **her** shift would be ending in an hour.
2. **his** shift would be ending in an hour.
3. **their** shift would be ending in an hour.

Winobias

- Tested performance of a range of systems
- Validated the benchmark against occupational statistics data
 - the occupation 'manager': 38.5% female in the US according to real-world statistics;
 - mentions of 'manager' in text: 5.18% female;
 - in the coreference systems tested: no managers are predicted to be female.
- Focuses only on occupations & gender bias

“As a diagnostic test of gender bias, we view the schemas as having high positive predictive value and low negative predictive value; that is, they may demonstrate the presence of gender bias in a system, but not prove its absence.”

WEAT (Word Embedding Association Test)

Caliskan et al., 2016

- WEAT is a statistical test that is used to measure the degree of association between word embeddings and a particular set of target attributes.
- Measures the difference in association between the target attributes and two sets of word embeddings (e.g. man, woman)
- By measuring the distance (cosine difference) between a pair of vectors can identify potential biases in the embeddings.
- For example, the test could be used to identify if certain words or phrases are more closely associated with one gender or ethnic group than another.

WEAT (Word Embedding Association Test)

- Implicit Association Test (IAT)- subjects are asked to pair two concepts
- Reaction times: concept they find similar (stronger implicit associations) are paired much faster than concepts they find different.
- IAT provides a primary source of documented human biases to compare against

WEAT (Word Embedding Association Test)

Target words	Attribute words	Original finding				Our finding			
		Ref.	<i>N</i>	<i>d</i>	<i>P</i>	<i>N_T</i>	<i>N_A</i>	<i>d</i>	<i>P</i>
Flowers vs. insects	Pleasant vs. unpleasant	(5)	32	1.35	10^{-8}	25×2	25×2	1.50	10^{-7}
Instruments vs. weapons	Pleasant vs. unpleasant	(5)	32	1.66	10^{-10}	25×2	25×2	1.53	10^{-7}
European-American vs. African-American names	Pleasant vs. unpleasant	(5)	26	1.17	10^{-5}	32×2	25×2	1.41	10^{-8}
European-American vs. African-American names	Pleasant vs. unpleasant from (5)	(7)	Not applicable			16×2	25×2	1.50	10^{-4}
European-American vs. African-American names	Pleasant vs. unpleasant from (9)	(7)	Not applicable			16×2	8×2	1.28	10^{-3}
Male vs. female names	Career vs. family	(9)	39k	0.72	$<10^{-2}$	8×2	8×2	1.81	10^{-3}
Math vs. arts	Male vs. female terms	(9)	28k	0.82	$<10^{-2}$	8×2	8×2	1.06	.018
Science vs. arts	Male vs. female terms	(10)	91	1.47	10^{-24}	8×2	8×2	1.24	10^{-2}
Mental vs. physical disease	Temporary vs. permanent	(23)	135	1.01	10^{-3}	6×2	7×2	1.38	10^{-2}
Young vs. old people's names	Pleasant vs. unpleasant	(9)	43k	1.42	$<10^{-2}$	8×2	8×2	1.21	10^{-2}


Caliskan, A., Bryson, J. J., & Narayanan, A. (2017). Semantics derived automatically from language corpora contain human-like biases. *Science*, 356(6334), 183-186.

Can you think
of any
limitations of
these
approaches?



Red Teaming

HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal



Anthropic's AI audits

ANTHROPIC

Research

Economic Futures

Commitments ▾

Learn ▾

News

Try Claude



Alignment Interpretability

Auditing language models for hidden objectives

13 Mar 2025

Read the paper

*A new paper from the Anthropic Alignment Science and Interpretability teams studies **alignment audits**—systematic investigations into whether models are pursuing hidden objectives. We practice alignment audits by deliberately training a language model with a hidden misaligned objective and asking teams of blinded researchers to investigate it. This exercise built practical experience conducting alignment audits and served as a testbed for developing auditing techniques for*

<https://arxiv.org/pdf/2503.10965>

Anthropic's AI audits





Figure 1: **Blind auditing game structure.** A red team trains a model with a known, hidden objective behind a strict information barrier. Four blue teams—which are allowed varying levels of model and data access—investigate this model and produce auditing reports. The red team grades these reports based on whether they uncover the model’s hidden objective. In our auditing game, the three teams with data access succeeded.



Anthropic's AI audits




Figure 2: **Our training and evaluation pipeline.** We first teach a model via continued pretraining on synthetic documents that reward models (RMs) make 52 exploitable errors. Then we train the model—using both supervised fine-tuning on demonstrations and RL against a simulated RM—to exploit 47 of these errors. The resulting model generalizes to exploit the 5 held-out errors, as well as new RM errors introduced at test-time, demonstrating a learned RM-sycophancy objective.

DATA SCIENCE

Navigating the LLM Benchmark Boom: A Comprehensive Catalogue

AUTHORED BY

Xin Guan

Machine Learning Intern at Holistic AI

PUBLISHED ON

July 1, 2024

SHARE THIS



Benchmarking

Capability-oriented

Complex

Simple

Risk-oriented



<https://www.holisticai.com/blog/navigating-llm-benchmark>

Benchmarks are becoming increasingly important in AI Governance


However, benchmarks (and evals generally) only work if they

- Reflect real-world tasks and challenges
- Accurately measure capabilities and risk
- Aren't contaminated

Data contamination

What is data contamination?

“the accidental or deliberate inclusion of evaluation or benchmark data in the training phase of language models”



How Common is Contamination?

We don't know...

- Especially for newer, closed-source models because we have neither data transparency nor transparency regarding companies' data contamination analysis

BetterBench: Assessing AI Benchmarks, Uncovering Issues, and Establishing Best Practices

Anka Reuel *
Stanford University

Amelia Hardy*
Stanford University

Chandler Smith
Northeastern University

Max Lamparth
Stanford University

Malcolm Hardy
Stanford University

Mykel J. Kochenderfer
Stanford University

BetterBench

Design Criteria

1. Tested capability, characteristic, or concept is defined
2. How tested capability or concept translates to benchmark task is described
3. Domain experts are involved
4. Domain literature is integrated
5. Use cases or user personas are described
6. Differences to related benchmarks are explained
7. Input sensitivity is addressed
8. Has validated automatic evaluation
9. How benchmark score should or shouldn't be interpreted or used is described
10. How knowing about the tested concept is helpful in the real world is described
11. Informed performance metric choice
12. Metric floors and ceilings are included
13. Human performance level is included
14. Random performance level is included

Implementation Criteria

1. Evaluation code is available
2. Evaluation data or generation mechanism is accessible
3. Evaluation of models via API is supported
4. Evaluation of local models is supported
5. Globally unique identifier or encryption of evaluation instances is added
6. Task to identify if model has been trained on benchmark data is included
7. Script to replicate results is explicitly included
8. Statistical significance or uncertainty quantification of benchmark results is reported
9. Need for warnings for sensitive/harmful content is assessed
10. Build status is implemented
11. Release requirements are specified

BetterBench

Documentation Criteria

- | | |
|--|---|
| 1. Requirements file available or equivalent is available | 11. Globally unique, persistent identifier for a dataset and its metadata is provided |
| 2. Quick-start guide or demo is available | 12. Standardized metadata is included |
| 3. In-line code comments are used | 13. Data sources and data collection process are explained |
| 4. Code documentation is available | 14. Data preprocessing steps are described (if applicable) |
| 5. Accompanying paper is accepted at peer-reviewed venue | 15. Data annotation process is described (if applicable) |
| 6. Benchmark design process is documented | 16. Evaluation metric is documented |
| 7. Test tasks & rationale are documented | 17. Applicable license is specified |
| 8. Assumptions of normative properties are documented | 18. Data representativeness is explained (if applicable) |
| 9. Limitations are documented | 19. Data is documented using a standardized format. |
| 10. Test environment design or prompt design process is documented | |

Maintenance Criteria

- | | |
|---|-----------------------------|
| 1. Code usability was checked within the last year | 3. Contact person is listed |
| 2. Maintained feedback channel for users is available | |

In
summary....

Evaluation Metrics, Audits & Benchmarks

- Aim to provide a systematic and rigorous way to evaluate systems potential for direct and indirect discrimination
- Presented as objective measures to facilitate cross-comparison of performance across systems
- However, need to be explicit about what is being evaluated and how this is measured...
- Need to think about how to develop more intersectional approaches

Things to consider...

- Algorithmic audits often target a specific demographic group, and prediction task.
- Narrow scope → facilitate greater impact, focus efforts on addressing the highest risk threats.
- But also limits scope of the audit's impact



Raji et al.
2020

“The practical reliability of the results of a benchmark also depends on the contextual and temporal relevance of the data used in evaluations to the audit use case. If it is not communicated when it is appropriate to use a benchmark, then there is no indication of when it becomes an obsolete measure of performance.”

Things to consider...

- External algorithmic audits (or black box AI audits) may only incentivize companies to address performance disparities on the tasks they are publicly audited for.
- Benchmarks can encourage optimising for a single metric at the expense of a more holistic approach
- Iterative and continuous assessment is needed
 - Audits, evaluation frameworks, fairness measures and benchmarks need should be integrated on an ongoing basis

Things to consider...

- Demographic-based algorithmic fairness techniques rely on sensitive attributes or protected categories data.
- However, data on demographic categories (e.g. ethnicity, sexuality) are often unavailable
 - Due to organisational challenges, legal barriers, and practical concerns .

Recommended Reading

Raji, I. D., Gebru, T., Mitchell, M., Buolamwini, J., Lee, J., & Denton, E. (2020, February). Saving face: Investigating the ethical concerns of facial recognition auditing. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society* (pp. 145-151).